

Concurrency and Parallelism Project

Guilherme Fernandes

Vladyslav Mikytiv

60045

60735

ABSTRACT

This report provides the explanation and the implementation of parallelisation of the algorithm of equalization of images. This is not the final report it's just a middle checkpoint.

KEYWORDS

OpenMP Parallelism C++

1 CRITICAL CODE ANALYSIS AND PARALLELIZATION

In this section we will analyse the behaviour of our code and identify the critical hotspots.

1.1 The first steps

At first we ran the code without doing any modifications and used the profiler in order to identify the hot spots. The zones of the code that took a long time to perform the computation and in overall slow down the algorithm are: the **normalization** of the image, the **correction of color**, the **conversion to the greyscale** and some other not so computational heavy functions.

To reduce the impact on processing time caused by these functions within the algorithm, we integrated OpenMP directives into these specific parts of the code to accelerate computation. Prior to this optimization, we structured the code by creating separate functions for each step of the algorithm to maintain organization.

It's important to highlight the **thread management**. Since we are working with images we will have to do some simples calculus to determine the amount of work that each thread will be given. For that we must perform two calculations (one for RGB and one for the grey scale). For the RGB we will calculate $WIDTH * HEIGHT * 3$ and to obtain the `CHUNK_SIZE_RGB` we just divide $WIDTH * HEIGHT * 3$ by `N_THREADS`. For the grey scale it's similar but we don't multiply by 3 getting the `CHUNK_SIZE`. Now we will have the work that will be balanced between threads.

1.2 Function parallelization

The **normalization** function will be improved with the following code:

```
1 void normalize(//omitting for space) {
2     #pragma omp parallel for schedule(dynamic,
3         chunk_size_channels) num_threads(n_threads)
4     for (int i = 0; i < size_channels; i++)
5         uchar_image[i] = (unsigned char) (255 *
6             input_image_data[i]);
7 }
```

Listing 1: Normalization Function

The **grey scale conversion** will be improved in two ways. We will mix the `fill_histogram` function with this one in order to do everything in the same function. Besides that we also apply OpenMP directives.

```
1 void convertToGrayscale(//omitting for space) {
2     // filling the histogram with zeroes
3     #pragma omp parallel for reduction(+: histogram)
4     num_threads(n_threads)
5     for (int i = 0; i < size; i++){
6         auto r = uchar_image[3 * i];
7         auto g = uchar_image[3 * i + 1];
8         auto b = uchar_image[3 * i + 2];
9         gray_image[i] =
10             static_cast<unsigned char>
11             (0.21 * r + 0.71 * g + 0.07 * b);
12         histogram[gray_image[i]]++;
13     }
```

Listing 2: Grey Conversion Function

The function that **calculated the CDF** doesn't require any type of parallelization. It's a simple for loop that executes 256 iterations every time. Another function that we managed to simplify was the `cdf_min_loop`. Since we minimum will always be on the first position we just return it and that's how we compute the minimum of the CDF.

The **correct_color_loop** function was parallelized with the following directives:

```
1 void correct_color_loop(//omitting for space) {
2     #pragma omp parallel for schedule(static,
3         chunk_size_channels) num_threads(n_threads)
4     for (int i = 0; i < size_channels; i++)
5         uchar_image[i] = correct_color
6         (cdf[uchar_image[i]], cdf_min);
7 }
```

Listing 3: Color Correction Function

The last function to be parallelized is **rescale**. And it was also changed accordingly with the OpenMP directives.

```
1 void rescale(//omitting for space) {
2     #pragma omp parallel for schedule(dynamic,
3         chunk_size_channels) num_threads(n_threads)
4     for (int i = 0; i < size_channels; i++)
5         output_image_data[i] = static_cast<float>
6         (uchar_image[i]) / 255.0f;
7 }
```

Listing 4: Rescale Function

2 METRIC ANALYSIS

Now, let's assess the impact of these changes on the runtime of our program. For that we will use **speed up** and **efficiency** and we will execute every executing multiple times in order to get the mean value of the execution times.

The code was executed on a computer with the following specifications: **Intel(R) Core(TM) i7-9750H CPU @ 2.60GHz** with a base speed of **2,59 GHz** and with **6 cores** and **12 logical processors**. The code was executed in a docker container.




Figure 1: Breakdown of the distribution of execution time across different number of threads usage

We can clearly observe that the execution time is always getting lower until it hits a certain point where it spikes again. The best NR_THREADS would be 11. Since the outlier is not as distant as with 12 threads it would be a better choice.

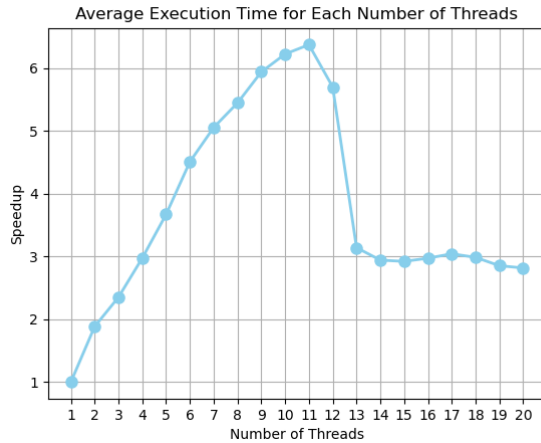


Figure 2: Speedup across different number of threads usage

In this graphic we can observe a very related behaviour to Figure 1 where the speed up keeps getting better until a certain point where it stops to increase and de-creases. Again we can observe that the best NR_THREADS would be 11.

In order to calculate the efficiency we will use the best result that we had: the one with 11 threads. For that we define S_{11} which

is the speed up with 11 threads and it will be equal to

$$S_{11} = 6.5$$

And the efficiency of the parallel implementation in Section 1 get's us an efficiency value of:

$$E_{11} = \frac{S_{11}}{11} = 59\%$$

3 FUTURE WORK

Now after applying the changes with OpenMP in our code we will enter in the realm of GPU computing. By using CUDA we will be able to do some computation in parallel in the GPU and send the results back to the CPU in order to accelerate even more our computation. We will have to do an analysis to see what is worth to parallelize and what is not.

4 CODE

All the code mentioned above can be consulted in <https://github.com/Gui28F/CP-Project> in the following commit ID 6a20ac8.